

Ch. 5. Neural Network and Deep Learning

Yongjae Yoo, Ph. D.
Assistant Professor
Department of Artificial Intelligence,
Hanyang University ERICA



So Far,

- We have learnt how your computer understand the data and make decisions.
- (In the simplest ways.)
- Based on the data, the machine can identify
 - The trends of data
 - Which group the new input belongs to

What If

- The question goes “deep,” i.e.,
 - Questions gets sophisticated – multiple times we should ask.
 - Questions have many possibilities
 - New point (input) has many dimensions
 - New point (input) has totally new one.
- Then, the basic approach goes impossible.
- In most cases in real life, we have such issues.

Examples

- Coke vs. Pepsi
 - So many ingredients, answers diverge.
- Conditional/confounded plans
 - “It depends on whether my friends’ opinions.”
 - We eat 빵. Turkish eat Ekmek. Chinese eat Mian Bao.
 - What will you do after this class?
 - Etc.

Dealing With Many Questions

- Nested If statements

```
x = 41

if x > 10:
    print("Above ten,")
    if x > 20:
        print("and also above 20!")
    else:
        print("but not above 20.")
```

- With and/or/not (binary conditional operators)

Binary Conditional Operators

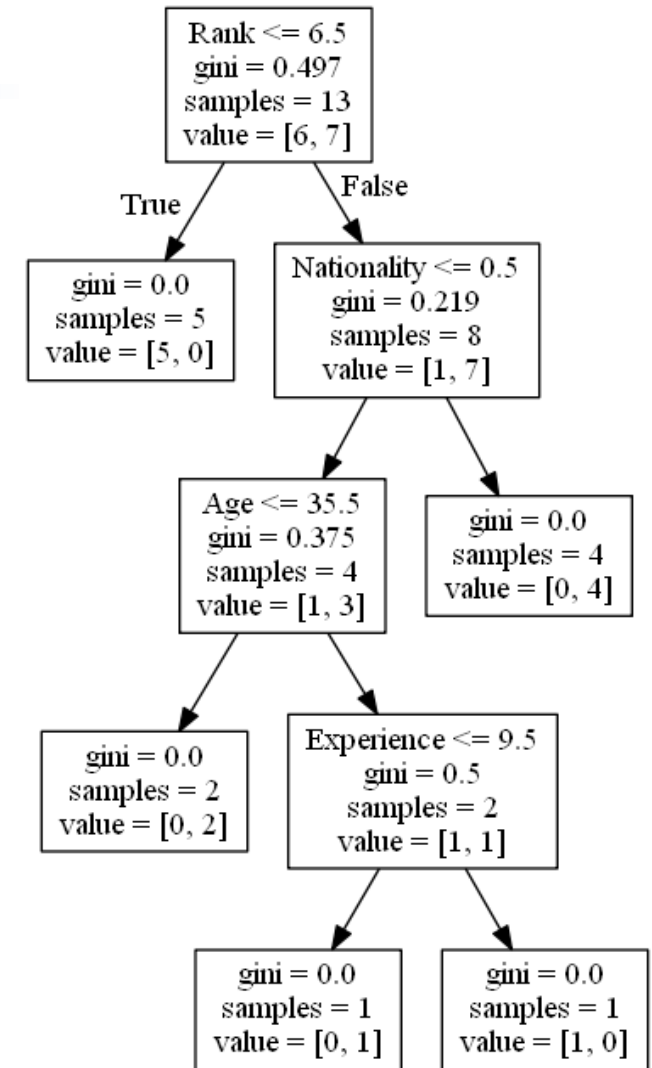
- And, and OR

```
a = 200
b = 33
c = 500
if a > b and c > a:
    print("Both conditions are True")
```

```
a = 200
b = 33
c = 500
if a > b or a > c:
    print("At least one of the conditions is True")
```

Decision Tree

- Dealing with multiple conditions in a systematic way.
- A series of true-false questions, and the answers will be decided by the sequential inputs.



Decision Tree

- Example:
 - “I am trying to decide if I should go to a comedy show or not.”
- Fortunately, I know the comedian’s experience, age, nationality, rank, and my preference (Yay or Nay).

| Age | Experience | Rank | Nationality | Go |
|-----|------------|------|-------------|-----|
| 36 | 10 | 9 | UK | NO |
| 42 | 12 | 4 | USA | NO |
| 23 | 4 | 6 | N | NO |
| 52 | 4 | 4 | USA | NO |
| 43 | 21 | 8 | USA | YES |
| 44 | 14 | 5 | UK | NO |
| 66 | 3 | 7 | N | YES |
| 35 | 14 | 9 | UK | YES |
| 52 | 13 | 7 | N | YES |
| 35 | 5 | 9 | N | YES |
| 24 | 3 | 5 | USA | NO |
| 18 | 3 | 7 | UK | YES |
| 45 | 9 | 9 | UK | YES |

Python Codes

- We will use a new tool, Pandas in this example.

```
import pandas
```

```
df = pandas.read_csv("data.csv")
```

```
print(df)
```

- Then, we need to convert the text data into numbers.
- We will use Pandas' map() method.

```
{'UK': 0, 'USA': 1, 'N': 2}
```

Python Codes

- Mapping texts into numbers:

```
d = {'UK': 0, 'USA': 1, 'N': 2}
df['Nationality'] = df['Nationality'].map(d)
d = {'YES': 1, 'NO': 0}
df['Go'] = df['Go'].map(d)

print(df)
```

- Then,

Python Codes

- Then, we will extract “*features*” to make decisions

```
features =  
['Age', 'Experience', 'Rank', 'Nationality']  
  
X = df[features]  
y = df['Go']  
  
print(X)  
print(y)
```

Final Decision Tree Here:

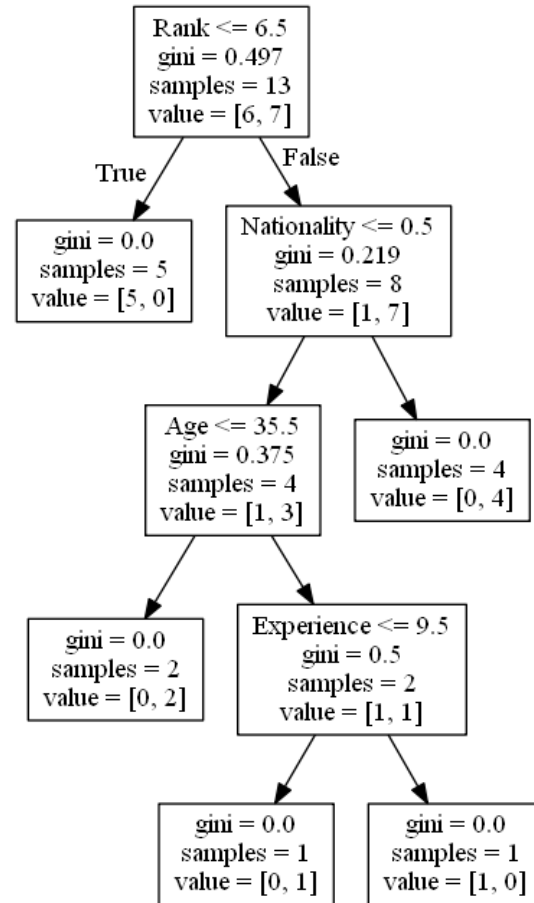
```
import pandas
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt

df = pandas.read_csv("data.csv")
d = {'UK': 0, 'USA': 1, 'N': 2}
df['Nationality'] = df['Nationality'].map(d)
d = {'YES': 1, 'NO': 0}
df['Go'] = df['Go'].map(d)

features = ['Age', 'Experience', 'Rank', 'Nationality']
X = df[features]
y = df['Go']


dtree = DecisionTreeClassifier()
dtree = dtree.fit(X, y)
tree.plot_tree(dtree, feature_names=features)
```

Results




Explaining Each Level

- It shows the criteria of decision on each
- First, see “rank.”
 - Rank ≤ 6.5 means that every comedian with a rank of 6.5 or lower will follow the **True** arrow (to the left), and the rest will follow the **False** arrow (to the right).
- Gini is the split quality – between 0 – 0.5
 - $\text{Gini} = 1 - (x/n)^2 - (y/n)^2$
Where x is the number of positive answers("GO") , n is the number of samples, and y is the number of negative answers ("NO"), which gives us this calculation: $1 - (7 / 13)^2 - (6 / 13)^2 = 0.497$




Explaining Each Level

- **samples = 13** means that there are 13 comedians left at this point in the decision, which is all of them since this is the first step.
- **value = [6, 7]** means that of these 13 comedians, 6 will get a "NO", and 7 will get a "GO".



2nd Level

- True: five comedians will end here.
- (five samples, and five “NO”s.
- False:
 - Then the next way is checking nationality.
 - Among 8 samples, 1 No and 7 Yeses.
 - Gini would be 0.219.



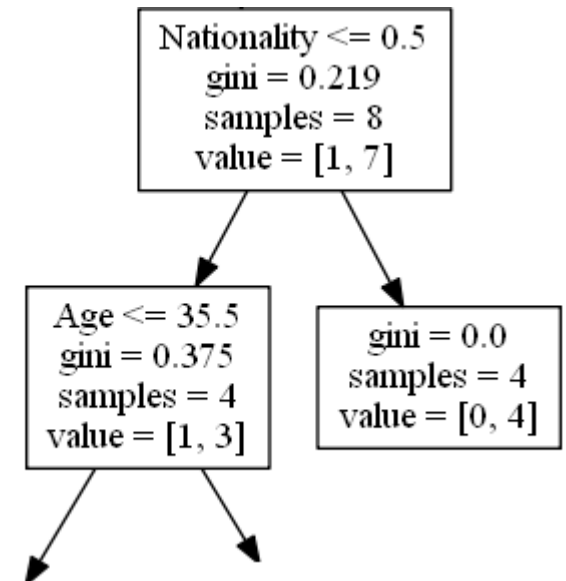
3rd Level

- True - 4 Comedians Continue:

- **Age** ≤ 35.5 means that comedians at the age of 35.5 or younger will follow the arrow to the left, and the rest will follow the arrow to the right.
- **gini** = 0.375 means that about 37.5% of the samples would go in one direction.
- **samples** = 4 means that there are 4 comedians left in this branch (4 comedians from the UK).
- **value** = [1, 3] means that of these 4 comedians, 1 will get a "NO" and 3 will get a "GO".

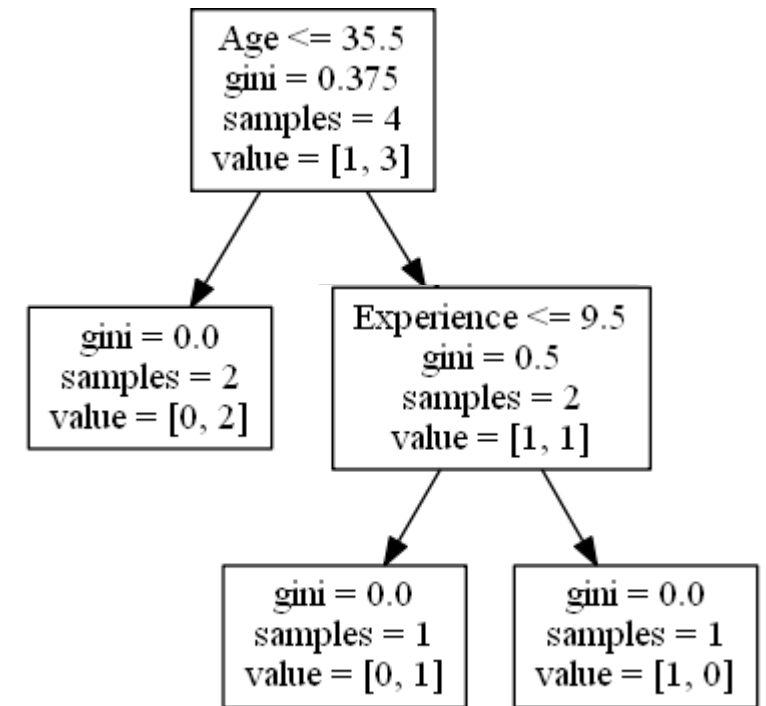
- False - 4 Comedians End Here:

- **gini** = 0.0 means all of the samples got the same result.
- **samples** = 4 means that there are 4 comedians left in this branch (4 comedians not from the UK).
- **value** = [0, 4] means that of these 4 comedians, 0 will get a "NO" and 4 will get a "GO".



4th and 5th levels

- In 4th level, it checks age.
 - If younger than 35.5, 2 samples will end with “Yes.”
 - Otherwise, it will check experience.
If one’s experience is shorter than 9.5 years,
it ends up to “No.”
Otherwise, goes to “Yes.”






So... Is The Decision Tree Always Works?

- Testing with different samples:
 - `print(dtree.predict([[40, 10, 7, 1]]))`
 - `print(dtree.predict([[40, 10, 6, 1]]))`
- You will see that the Decision Tree gives you different results, even you put in the same data when you try multiple times.
- That is because the Decision Tree is based on the probability of an outcome, and the answer will vary.

Curse of Dimensionality

- If we set multiple domain(s) – variables,
- The possibility and the number of calculation explodes.
- In your cases of



Anyhow Let's Do This



- Try the example.
- If you have any good data, you can try that with your data for your final report.

Dealing with Complexity



- Once again, computer can deal with yes/no questions.
- In other words, it can answer “data with linearity.”
- In theory, we can deal with MANY linear combinations of questions if we have enough memory. That is so-called Linear score function.
- To resolve this, we will let computer think as like we do.
- Neural network!

Analogy to (Physiologic) Neurons

- Physical neural networks vs. computational neural networks.




Looking into Single Neuron



Neuron as The Basic Building Block


- A neuron takes inputs, does some math with them, and produces one output.
- 2-input neuron:



In a Neuron:

Inputs

Output



3 things are happening here. First, each input is multiplied by a weight: ■

$$x_1 \rightarrow x_1 * w_1$$

$$x_2 \rightarrow x_2 * w_2$$

Next, all the weighted inputs are added together with a bias b : ■

$$(x_1 * w_1) + (x_2 * w_2) + b$$

Finally, the sum is passed through an activation function: ■

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

Activation function


- Works as a switch – it cuts out conditions we do not want.

$$f = W_2 \max(0, W_1 x)$$

Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Rectified Linear Unit (ReLU) often used as Default

Activation Function

- Consider a sigmoid function
 - Converts $-\infty$ to $+\infty$ to $(0, 1)$



Example

- Assume we have a 2-input neuron that uses the sigmoid activation function and has the following parameters:

$$w = [0, 1]$$

$$b = 4$$

- $w = [0, 1]$: $w_1 = 0, w_2 = 1$.
- Now give an input $x = [2, 3]$. Applying a dot product:

$$\begin{aligned}(w \cdot x) + b &= ((w_1 * x_1) + (w_2 * x_2)) + b \\ &= 0 * 2 + 1 * 3 + 4 \\ &= 7\end{aligned}$$

$$y = f(w \cdot x + b) = f(7) = \boxed{0.999}$$

A Single Neuron's Output

- Gives you a single output here – 0.999, as a “feedforward” result.
- Feedforward process means:
passing inputs forward to get an output

Now, Coding a Neuron Here:

- Here just we implemented example converted to code:

```
import numpy as np

def sigmoid(x): # Our activation function:  $f(x) = 1 / (1 + e^{-x})$ 
    return 1 / (1 + np.exp(-x))


class Neuron:
    def __init__(self, weights, bias):
        self.weights = weights
        self.bias = bias

    def feedforward(self, inputs):
        # Weight inputs, add bias, then use the activation function
        total = np.dot(self.weights, inputs) + self.bias
        return sigmoid(total)

weights = np.array([0, 1]) # w1 = 0, w2 = 1
bias = 4 # b = 4
n = Neuron(weights, bias)
x = np.array([2, 3]) # x1 = 2, x2 = 3
print(n.feedforward(x)) # 0.9990889488055994
```


So, What's The Difference?

- Neural network can be concatenated, like:



- This network has 2 inputs, a hidden layer with 2 neurons h_1 and h_2 , and an output layer with 1 neuron (o_1).
- Notice that the inputs for o_1 are the outputs from h_1 and h_2 - that's what makes this a network.

Feedforwarding the Network

- A neural network can have any number of layers with any number of neurons in those layers.
- The basic idea stays the same: feed the input(s) forward through the neurons in the network to get the output(s) at the end.

What happens if we pass in the input $x = [2, 3]$?

$$\begin{aligned}h_1 = h_2 &= f(w \cdot x + b) \\&= f((0 * 2) + (1 * 3) + 0) \\&= f(3) \\&= 0.9526\end{aligned}$$

$$\begin{aligned}o_1 &= f(w \cdot [h_1, h_2] + b) \\&= f((0 * h_1) + (1 * h_2) + 0) \\&= f(0.9526) \\&= \boxed{0.7216}\end{aligned}$$

Let's Implement a 2-Layer Network

```
import numpy as np # ... code from previous section here

class OurNeuralNetwork:
    ''' A neural network with:
        - 2 inputs
        - a hidden layer with 2 neurons (h1, h2)
        - an output layer with 1 neuron (o1)
        - Each neuron has the same weights and bias:
            - w = [0, 1]
            - b = 0 '''

    def __init__(self):
        weights = np.array([0, 1])
        bias = 0

        # The Neuron class here is from the previous section
        self.h1 = Neuron(weights, bias)
        self.h2 = Neuron(weights, bias)
        self.o1 = Neuron(weights, bias)

    def feedforward(self, x):
        out_h1 = self.h1.feedforward(x)
        out_h2 = self.h2.feedforward(x)

        # The inputs for o1 are the outputs from h1 and h2
        out_o1 = self.o1.feedforward(np.array([out_h1, out_h2]))
        return out_o1

network = OurNeuralNetwork()


x = np.array([2, 3])
print(network.feedforward(x)) # 0.7216325609518421
```

So, More about Training

- Assume we have the following measurements:

| Name | Weight (lb) | Height (in) | Gender |
|---------|-------------|-------------|--------|
| Alice | 133 | 65 | F |
| Bob | 160 | 72 | M |
| Charlie | 152 | 70 | M |
| Diana | 120 | 60 | F |

- Let's train our network to predict someone's gender given their weight and height:



So, More about Training

- We'll represent Male with a 0 and Female with a 1, and we'll also shift the data to make it easier to use:

| Name | Weight (minus 135) | Height (minus 66) | Gender |
|---------|--------------------|-------------------|--------|
| Alice | -2 | -1 | 1 |
| Bob | 25 | 6 | 0 |
| Charlie | 17 | 4 | 0 |
| Diana | -15 | -6 | 1 |

*I arbitrarily chose the shift amounts (135 and 66) to make the numbers look nice.
Normally, you'd shift by the mean.*

Understanding Loss

- Before we train our network, we first need a way to quantify how “good” it’s doing so that it can try to do “better”.
- That’s what the loss is.

We’ll use the **mean squared error** (MSE) loss:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_{\text{true}} - y_{\text{pred}})^2$$

the **squared error**.

- n is the number of samples, which is 4 (Alice, Bob, Charlie, Diana).
- y represents the variable being predicted, which is Gender.
- y_{true} is the *true* value of the variable (the “correct answer”). For example, y_{true} for Alice would be 1 (Female).
- y_{pred} is the *predicted* value of the variable. It’s whatever our network outputs.

Understanding Loss

- The Keys:
 - Better predictions = Lower loss.
 - Training a network = trying to minimize its loss.

Example

An Example Loss Calculation

Let's say our network always outputs 0 - in other words, it's confident all humans are Male 🤔. What would our loss be?

| Name | y_{true} | y_{pred} | $(y_{true} - y_{pred})^2$ |
|---------|------------|------------|---------------------------|
| Alice | 1 | 0 | 1 |
| Bob | 0 | 0 | 0 |
| Charlie | 0 | 0 | 0 |
| Diana | 1 | 0 | 1 |

$$\text{MSE} = \frac{1}{4}(1 + 0 + 0 + 1) = \boxed{0.5}$$

Here MSE Loss Code

- We can convert the procedure as code.

```
import numpy as np


def mse_loss(y_true, y_pred):
    # y_true and y_pred are numpy arrays of the same length.
    return ((y_true - y_pred) ** 2).mean()

y_true = np.array([1, 0, 0, 1])
y_pred = np.array([0, 0, 0, 0])

print(mse_loss(y_true, y_pred)) # 0.5
```

Getting A Little Bit More Difficult

- So-called Backprop
- We now have a clear goal: minimize the loss of the neural network.
- We will skip the details of math here, but general way of updating the parameters here.



BackProp - Mathematics

To start, let's rewrite the partial derivative in terms of $\frac{\partial y_{pred}}{\partial w_1}$ instead:

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y_{pred}} * \frac{\partial y_{pred}}{\partial w_1}$$

This works because of the [Chain Rule](#).

We can calculate $\frac{\partial L}{\partial y_{pred}}$ because we computed $L = (1 - y_{pred})^2$ above:

$$\frac{\partial L}{\partial y_{pred}} = \frac{\partial (1 - y_{pred})^2}{\partial y_{pred}} = \boxed{-2(1 - y_{pred})}$$

Now, let's figure out what to do with $\frac{\partial y_{pred}}{\partial w_1}$. Just like before, let h_1, h_2, o_1 be the outputs of the neurons they represent. Then

$$y_{pred} = o_1 = f(w_5 h_1 + w_6 h_2 + b_3)$$

f is the sigmoid activation function, remember?

BackProp - Mathematics

Since w_1 only affects h_1 (not h_2), we can write

$$\frac{\partial y_{pred}}{\partial w_1} = \frac{\partial y_{pred}}{\partial h_1} * \frac{\partial h_1}{\partial w_1}$$
$$\frac{\partial y_{pred}}{\partial h_1} = \boxed{w_5 * f'(w_5 h_1 + w_6 h_2 + b_3)}$$

More Chain Rule.

We do the same thing for $\frac{\partial h_1}{\partial w_1}$:

$$h_1 = f(w_1 x_1 + w_2 x_2 + b_1)$$
$$\frac{\partial h_1}{\partial w_1} = \boxed{x_1 * f'(w_1 x_1 + w_2 x_2 + b_1)}$$

You guessed it, Chain Rule.

x_1 here is weight, and x_2 is height. This is the second time we've seen $f'(x)$ (the derivative of the sigmoid function) now! Let's derive it:

$$f(x) = \frac{1}{1 + e^{-x}}$$
$$f'(x) = \frac{e^{-x}}{(1 + e^{-x})^2} = f(x) * (1 - f(x))$$

BackProp - Mathematics

We're done! We've managed to break down $\frac{\partial L}{\partial w_1}$ into several parts we can calculate:

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y_{pred}} * \frac{\partial y_{pred}}{\partial h_1} * \frac{\partial h_1}{\partial w_1}$$

We will use this derived equation
For actual calculation 😊

This system of calculating partial derivatives by working backwards is known as **backpropagation**, or "backprop".

Training: Stochastic Gradient Descent

- We'll use an optimization algorithm called stochastic gradient descent (SGD) that tells us how to change our weights and biases to minimize loss.
- It's basically just this update equation:

$$w_1 \leftarrow w_1 - \eta \frac{\partial L}{\partial w_1}$$

- η : learning rate
 - If $\frac{\partial L}{\partial w_1}$ is positive, w_1 will decrease, which makes L decrease.
 - If $\frac{\partial L}{\partial w_1}$ is negative, w_1 will increase, which makes L decrease.


Just Jumping into Calculations

- Choose a single item (stochastic) and try to learn from data.
 1. Choose **one** sample from our dataset. This is what makes it *stochastic* gradient descent - we only operate on one sample at a time.
 2. Calculate all the partial derivatives of loss with respect to weights or biases (e.g. $\frac{\partial L}{\partial w_1}$, $\frac{\partial L}{\partial w_2}$, etc).
 3. Use the update equation to update each weight and bias.
 4. Go back to step 1.

Let's Make a Complete Implementation

It's *finally* time to implement a complete neural network:

| Name | Weight (minus 135) | Height (minus 66) | Gender |
|---------|--------------------|-------------------|--------|
| Alice | -2 | -1 | 1 |
| Bob | 25 | 6 | 0 |
| Charlie | 17 | 4 | 0 |
| Diana | -15 | -6 | 1 |



Codes (1)

```
import numpy as np

def sigmoid(x):
    # Sigmoid activation function:  $f(x) = 1 / (1 + e^{-x})$ 
    return 1 / (1 + np.exp(-x))

def deriv_sigmoid(x):
    # Derivative of sigmoid:  $f'(x) = f(x) * (1 - f(x))$ 
    fx = sigmoid(x)
    return fx * (1 - fx)

def mse_loss(y_true, y_pred):
    # y_true and y_pred are numpy arrays of the same length.
    return ((y_true - y_pred) ** 2).mean()
```

Codes (2)

```
class OurNeuralNetwork:
    ''' A neural network with:
    - 2 inputs, a hidden layer with 2 neurons (h1, h2), an output layer with 1 neuron (o1)
    '''
    def __init__(self):
        # Weights
        self.w1 = np.random.normal()
        self.w2 = np.random.normal()
        self.w3 = np.random.normal()
        self.w4 = np.random.normal()
        self.w5 = np.random.normal()
        self.w6 = np.random.normal()
        # Biases
        self.b1 = np.random.normal()
        self.b2 = np.random.normal()
        self.b3 = np.random.normal()

    def feedforward(self, x):
        # x is a numpy array with 2 elements.
        h1 = sigmoid(self.w1 * x[0] + self.w2 * x[1] + self.b1)
        h2 = sigmoid(self.w3 * x[0] + self.w4 * x[1] + self.b2)
        o1 = sigmoid(self.w5 * h1 + self.w6 * h2 + self.b3)
        return o1
```

Codes (3)

```
def train(self, data, all_y_trues):
    ''' - data is a (n x 2) numpy array, n = # of samples in the dataset.
        - all_y_trues is a numpy array with n elements.
        Elements in all_y_trues correspond to those in data. '''

    learn_rate = 0.1 epochs = 1000
    # number of times to loop through the entire dataset
    for epoch in range(epochs):
        for x, y_true in zip(data, all_y_trues):
            # --- Do a feedforward (we'll need these values later)
            sum_h1 = self.w1 * x[0] + self.w2 * x[1] + self.b1 h1 = sigmoid(sum_h1)
            sum_h2 = self.w3 * x[0] + self.w4 * x[1] + self.b2 h2 = sigmoid(sum_h2)
            sum_o1 = self.w5 * h1 + self.w6 * h2 + self.b3
            o1 = sigmoid(sum_o1) y_pred = o1

            # --- Calculate partial derivatives.
            # --- Naming: d_L_d_w1 represents "partial L / partial w1"
            d_L_d_ypred = -2 * (y_true - y_pred)
            # Neuron o1
            d_ypred_d_w5 = h1 * deriv_sigmoid(sum_o1)
            d_ypred_d_w6 = h2 * deriv_sigmoid(sum_o1)
            d_ypred_d_b3 = deriv_sigmoid(sum_o1)
            d_ypred_d_h1 = self.w5 * deriv_sigmoid(sum_o1)
            d_ypred_d_h2 = self.w6 * deriv_sigmoid(sum_o1)
```

```
# Neuron h1
d_h1_d_w1 = x[0] * deriv_sigmoid(sum_h1)
d_h1_d_w2 = x[1] * deriv_sigmoid(sum_h1)
d_h1_d_b1 = deriv_sigmoid(sum_h1)
# Neuron h2
d_h2_d_w3 = x[0] * deriv_sigmoid(sum_h2)
d_h2_d_w4 = x[1] * deriv_sigmoid(sum_h2)
d_h2_d_b2 = deriv_sigmoid(sum_h2)

# --- Update weights and biases
# Neuron h1
self.w1 -= learn_rate * d_L_d_ypred * d_ypred_d_h1 * d_h1_d_w1
self.w2 -= learn_rate * d_L_d_ypred * d_ypred_d_h1 * d_h1_d_w2
self.b1 -= learn_rate * d_L_d_ypred * d_ypred_d_h1 * d_h1_d_b1

# Neuron h2
self.w3 -= learn_rate * d_L_d_ypred * d_ypred_d_h2 * d_h2_d_w3
self.w4 -= learn_rate * d_L_d_ypred * d_ypred_d_h2 * d_h2_d_w4
self.b2 -= learn_rate * d_L_d_ypred * d_ypred_d_h2 * d_h2_d_b2

# Neuron o1
self.w5 -= learn_rate * d_L_d_ypred * d_ypred_d_w5
self.w6 -= learn_rate * d_L_d_ypred * d_ypred_d_w6
self.b3 -= learn_rate * d_L_d_ypred * d_ypred_d_b3


# --- Calculate total loss at the end of each epoch
if epoch % 10 == 0:
    y_preds = np.apply_along_axis(self.feedforward, 1, data)
    loss = mse_loss(all_y_trues, y_preds)
    print("Epoch %d loss: %.3f" % (epoch, loss))
```

Codes (4)

```
# Define dataset
data = np.array([ [-2, -1], # Alice [25, 6], # Bob [17, 4], # Charlie [-15, -6], # Diana ])
all_y_trues = np.array([ 1, # Alice 0, # Bob 0, # Charlie 1, # Diana ])

# Train our neural network!
network = OurNeuralNetwork()
network.train(data, all_y_trues)
```

- Training decreases the loss



Making Actual Prediction

```
# Make some predictions
emily = np.array([-7, -3]) # 128 pounds, 63 inches
frank = np.array([20, 2]) # 155 pounds, 68 inches
print("Emily: %.3f" % network.feedforward(emily)) # 0.951 - F
print("Frank: %.3f" % network.feedforward(frank)) # 0.039 - M
```

More Resource

- Victor Zhou's blog page on easy introduction to NN
- <https://victorzhou.com/blog/intro-to-neural-networks/#code-a-complete-neural-network>
- <https://replit.com/@vzhou842/An-Introduction-to-Neural-Networks>

Lab Session

- Follow the whole code and try to understand.
- Consider which of prediction you can make with this.
 - What is your major?

Dataset Treasure Island: Kaggle

- www.kaggle.com

